

Exercise 1

Task 1

Subtask A

term	doc freq.	postings list
beach	5	1, 2, 3, 4, 7
summer	3	1, 5, 6
holiday	4	5, 6, 8, 9
is	1	6
beer	1	10

Subtask B

Skip pointers added:

term	doc freq.	postings list
beach	5	1, 2, 3, 4, 7
summer	3	1, 5, 6
holiday	4	5, 6, 8, 9
is	1	6
beer	1	10

Example:

- Add skip pointer for postings list for the term beach from 1 to 4
- The query:
 - **beach AND holiday**
- We start at the postings lists with the entries **beach = 1** and **holiday = 5**
- So we can take the skip pointer at the term beach from 1 to 4, so we are at beach = 4
- Next we compare beach = 7 → Now we know that there is no match
- **The skip pointer saved us the comparisons of 1,2,3 for the term beach.**
 - Without skip pointers we need more comparisons, thus this query can be answered in a more efficient way with these skip pointers.

Task 2


Pseudo-Code:

1. Initialize currentChar = 0
2. Initialize lastStop = 0
3. Initialize resultList = []
4. Iterate over inputString at position currentChar with condition currentChar < inputString.length
 - a. If inputString[currentChar] is Whitespace add substring of inputString starting from index lastStop, ending at index currentChar to resultList. Set lastStop = currentChar+1
 - b. If currentChar equals inputString.length – 1 add substring of inputString starting from index lastStop, ending at index currentChar+1 to resultList.
 - c. currentChar++ and continue iteration at 4
5. return resultList

Java-Code implementation:

```
1. public class Tokenizer {
2.     public static String example = "Mr. Nopper moved from Backnang to Stuttgart";
3.
4.     static List<String> tokenize(String input) {
5.         List<String> result = new ArrayList();
6.
7.         int lastStop = 0;
8.
9.         for(int currentChar=0; currentChar<input.length(); currentChar++) {
10.            if(Character.isWhitespace(input.charAt(currentChar))) {
11.                result.add(input.substring(lastStop, currentChar));
12.                lastStop = currentChar+1;
13.            }
14.            if(currentChar == input.length()-1) {
15.                result.add(input.substring(lastStop, currentChar+1));
16.            }
17.        }
18.        return result;
19.    }
20.
21.    public static void main(String[] args) {
22.        List<String> test = tokenize(example);
23.        test.forEach(t -> System.out.print(t+ "|"));
24.    }
25. }
```

Output:



The presented algorithm has in general a linear runtime, because it only iterates the `inputString` once with the `for-Loop` and the loop variable `currentChar` is never changed/reset inside the loop. Therefore the `for-loop` is executing exactly n -times for an `inputString` of length n . All other operations (variable initialization) have a runtime of $O(1)$. This results in a $O(n)$ runtime for the algorithm.

Task 3

$i \rightarrow$ $j \downarrow$ 0	H 1	u 2	a 3	s 4
H 1	0 stay	1 delete	2 delete	3 delete
a 2	1 insert	1 replace	1 stay	2 delete
u 3	2 insert	1 stay	1 transpose	2 delete, replace
s 4	3 insert	2 insert	2 insert, replace	1 stay

Operations		add	condition
stay	$D_{i-1, j-1}$	0	if $u_i = v_j$
replace	$D_{i-1, j-1}$	1	
insert	$D_{i, j-1}$	1	
delete	$D_{i-1, j}$	1	
transpose	$D_{i-2, j-2}$	1	$u_i = v_{j-1} \ \&\& \ u_{i-1} = v_j$

Explanation

- The Levenshtein distance between „Haus“ and „Huas“ is 2, because we need to perform two times the replace operation, replacing a with u and the u with a.
- The Damerau-Levenshtein distance between „Haus“ and „Huas“ is 1, because we need to perform one transpose operation, which switches the letters u and a.
- The operations that produced the value in the matrix are described in the table. If there is more than one possible operation these are mentioned in a comma separated list.
- The final distance can be retrieved from the bottom right corner of the matrix.

We can backtrack the steps starting from the bottom right corner:

- Stay → step -1, -1 in matrix
- Transpose → step -2,-2 in matrix
- Stay → We reached the the top left corner of the matrix

Task 4

Export Tweets to disk → createTweetDB.py

```
import pandas as pd
import numpy as np
# this script exports all tweets as .txt files to disk. The filename is the tweet handle

nrows = 0
folder = "E:/tweets/"

# read csv
print("reading csv...")
if(nrows>0):
    # only read nrows number of tweets
    data = pd.read_csv('twitter.csv', sep="\t",
                      names=["handle", "userid", "username", "tweet"],
                      dtype={"handle": np.int64, "userid": str, "username": str, "tweet": str},
                      header=None, nrows=nrows,
                      quoting=3)
else:
    # read all tweets
    data = pd.read_csv('twitter.csv', sep="\t",
                      names=["handle", "userid", "username", "tweet"],
                      dtype={"handle": np.int64, "userid": str, "username": str, "tweet": str},
                      header=None,
                      quoting=3)

# export tweets to txt files where filename is equal to the tweet handle
for item in data.iterrows():
    with open(folder+str(item.handle)+'.txt', 'w', encoding='utf-8') as f:
        f.write(item.tweet)
```

Build index file → createIndex.py

```
import csv
from datetime import datetime
from itertools import chain
import numpy as np
import nltk
import pandas as pd
import pickle

# This script builds a non-positional inverted index. Postings lists are stored as .csv files on disk.
# The index is stored as pickled python object on disk and can be loaded by the queryEngine later.
# Stopwords of the 8 most common languages of the tweets are filtered, as well as words with length = 1.

# The index is a dictionary with the term as key and the values docFrequency and postings
# docFrequency is an integer which counts how often the term occurs in all documents
# postings is an integer with the file ID for the postings list.
# E.g. postings=100 means, the posting list is stored in the file p100.csv

# config
indexName = datetime.today().strftime('%Y%m%d-%H_%M_%S') + "_index.pickle"
nrows = 0

# token normalization
def normalize(line):
    # this function removes special characters, newlines and tab from tweets, tokenizes the text and sets all terms
    # to lowercase
    text = line.tweet
    if (type(text) != str):
        return ""

    for ch in ['[NEWLINE]', '[TAB]', '#', '\\', '\t', '*', '_', '{', '}', '[', ']', '(', ')', '>', '+', '-', '.', '!', '?', '$', '\'', '"', '/']:
        if ch in text:
            text = text.replace(ch, " ")
    return text.lower().split()

# index function
def index(filename):
    # read csv
    print("reading csv...")
    if (nrows > 0):
        data = pd.read_csv(filename, sep="\t",
                          names=["handle", "userid", "username", "tweet"],
                          dtype={"handle": np.int64, "userid": str, "username": str, "tweet": str},
                          header=None,
                          nrows=nrows,
                          quoting=3)
    else:
        data = pd.read_csv(filename, sep="\t",
                          names=["handle", "userid", "username", "tweet"],
                          dtype={"handle": np.int64, "userid": str, "username": str, "tweet": str},
                          header=None,
                          quoting=3)
```

```

# get stopwords
nltk.download('stopwords')
languages = ['english', 'german', 'spanish', 'portuguese', 'italian', 'french', 'turkish', 'dutch']
stopwords = dict.fromkeys([i for i in chain.from_iterable([nltk.corpus.stopwords.words(l) for l in languages])])

# gather data
progress = 0;
print("building inverted index ...")
results = {}
for item in data.iloc:

    progress = progress + 1;
    if (progress % 10000 == 0):
        print("processing line " + str(progress) + " ...")

    # tokenize tweet
    terms = normalize(item)

    for term in terms:

        # ignore stopwords and single characters
        if term in stopwords or len(term) < 2:
            continue

        if term in results:
            entry = results[term]
            entry["docFreq"] = entry["docFreq"] + 1
            entry["postings"].add(item.handle)
        else:
            results[term] = {"docFreq": 1, "postings": set([item.handle])}

# save postings lists as file
print("create postings lists files...")
fileId = 0

for term in results:
    # for each term write postings lists to csv file on disk
    entry = results[term]
    fileName = str(fileId) + '.csv'
    file = open('E:/postings_lists/p' + fileName, 'w', newline='')
    writer = csv.writer(file)
    postings = sorted(list(entry['postings']))
    writer.writerows([[handle] for handle in postings])
    entry['postings'] = fileId
    fileId = fileId + 1

    if fileId % 10000 == 0:
        print("writing file " + str(fileId))

# pickle index as file on disk
print("pickling results...")
f = open(indexName, 'wb')
pickle.dump(results, f)
f.close()
print("done.")
return results

# run
indexData = index("twitter.csv")

```


Query engine → queryEngine.py

```
import pickle
from pathlib import Path

import pandas as pd

# config
indexName = "20211108-14_03_27_index.pickle"
postings_dir = "E:/postings_lists/"
tweets_dir = "E:/tweets/"
showResults = 0
index = None

# functions
def query(term):
    if term in index:
        # read csv file from disk and return values as list
        postings = pd.read_csv(postings_dir + 'p' + str(index[term]['postings']) + '.csv')
        return postings.iloc[:, 0].values.tolist()
    else:
        # term not in index, return empty list
        return []

def queryAND(term1, term2):
    # get postings lists
    list1 = iter(query(term1))
    list2 = iter(query(term2))

    # intersection algorithm:

    # get start values smallValue, bigValue not sorted yet!
    smallV = next(list1, None)
    bigV = next(list2, None)

    # if one term has no results, AND query has no results
    if smallV == None or bigV == None:
        return []

    # init conditions → list1 with smallV tries to catchup to list2 to bigV
    if (smallV > bigV):
        temp = smallV
        temp2 = list1
        smallV = bigV
        list1 = list2
        bigV = temp
        list2 = temp2

    matches = []
    # intersect lists
    while (smallV <= bigV):
        # get value
        smallV = next(list1, None)

        # one list reached the end, break and return
        if (smallV == None):
            break

        # match found!
        if (smallV == bigV):
            matches.append(smallV)

        # smallValue passed bigValue => swap values and lists
        if (smallV > bigV):
            temp = smallV
            temp2 = list1
            smallV = bigV
            list1 = list2
            bigV = temp
            list2 = temp2

    return matches
```

```

def displayResults(postingslists, nOfResults=showResults):
    results = []

    # if results > nOfResults show only nOfResults many results
    maxIndex = len(postingslists)
    if nOfResults != 0 and nOfResults < maxIndex:
        maxIndex = nOfResults

    # if no results
    if maxIndex == 0:
        print("No results...")
        return results

    # iterate over results, open tweet.txt file from disk and print tweet in a line
    for i in range(0, maxIndex):
        file = tweets_dir + str(postingslists[i]) + ".txt"
        with open(file, 'r', encoding='utf-8') as f:
            results.append(str(postingslists[i]) + " : " + f.read())
    for i in range(0, len(results)):
        print(results[i])

    return results

# run the query loop:

# check if index file exists
results = Path(indexName)
if results.is_file():

    # unpickle index file
    print("file found, loading...")
    pickleFile = open(indexName, 'rb')
    index = pickle.load(pickleFile)

    # UI loop, let user enter queries
    while True:
        queryString = input("Enter your search query: ")

        # exit program with q
        if (queryString == "q"):
            print("Quit called. Goodbye, see you later alligator...")
            break

        # process query
        terms = queryString.lower().split()
        if (len(terms) == 1):
            # one term, use simple query
            results = query(terms[0])
        else:
            # two terms, use AND query
            results = queryAND(terms[0], terms[1])

        # display results
        resultDocuments = displayResults(results)

```

Results & Example query

```
Enter your search query: Minecraft Monday
1346239551342166016 : Starting with some Minecraft Monday for my week. Catch me on #Twitch streaming some Stoneblock2. Hope you are having a great Monday ♥
1363261466334306305 : #Minecraft #livestream Monday 10AM MT, come hang out on #twitch link in bio and will post at the start of my stream!
1371236823121399809 : Schlatt said he and Connor ran a business in another server, Quackity had canon PTSD of Techno from Minecraft Monday. So what we can conclude is that Schlatt and Connor
1398773118928408577 : people mentioning minecraft monday is bringing me back to watching vik and getting so pissed when techno used /top to beat him because I had no clue who techno was and
1398775542124486656 : People who got really really mad on mcc today probably wouldn't survive watching the previous mcc or minecraft monday.
1399408095055343617 : #MINECRAFT #MONDAY TIME!!
1412233747697598464 : Playing some MineCraft with my nephew. How's everyone's Monday been so far? #monday
1427345246912974870 : Running a little late, but Minecraft Monday will be going live imminently! Do tune in! It's not to be missed #Minecraft #monday #gameshow
1429221700978184195 : Thanks everyone who came out to my #Minecraft stream tonight; I had the absolute BEST time (despite the constant fear of unseen mobs and getting HORRENDOUSLY lost) and.
1429593841015967745 : Hey all! So as most know I recently got tested positive for Covid and delayed my streams. But now I'm feeling a little better. So we will be back at it tomorrow with ou
1429924186101895175 : What is up everyone I am no longer #DownWithTheSickness. So its time to get down with this weeks schedules...[NEWLINE][NEWLINE]Monday- CANCELED [NEWLINE][NEWLINE] Wednes
Enter your search query: Stuttgart silvester
1344936702369017859 : Wir leben in einem Deutschland, in dem an Silvester #Covidioten in #Stuttgart bis ins neue Jahr hinein demonstrieren dürfen. [NEWLINE][NEWLINE]Wir leben in einem Deuts
1345050982398251008 : Das Maß ist nach Silvester endgültig voll - Demokratie hat ihre Grenzen! Zeit jeden Corona-Leugner und #Querdenker0711 vor die Tore vor #Stuttgart zu treiben! Wir lasse
Enter your search query: Berlin uahn
1345800756021108738 : Heute in der Bahn wieder einen netten Mitmenschen getroffen. Jedoch war er mir nicht ganz geheuer, als er mir ein Gram Schore angeboten hat. Ich musste dankend ablehnen
1346476380565897219 : Ich fahre gerade leider (aus einem triftigen Grund) mit einer S-Bahn aus Berlin raus. Sagen wir so: Mehr Homeoffice wäre möglich und nötig.
1347972850863529987 : Deutschland Wirtschaft Läuft gut Unternehmer führen gut trotzdem wollen sich Politiker immer mehr einmischen jetzt Grün mit HomeOffice Pflicht stellt Grün Geräte Sicher
1349080173322706946 : Unterschied einer S-Bahn-Station: vor den 6 km: Inzidenz Berlin-Köpenick: 220. Schulen und Kitas zu. Hinter den 6 km: Inzidenz Dahme-Spreewald: 304. Abschlussklassen in
1350792695132921856 : Funktioniert bei euch die #CoronaApp noch? Hatte seit Weihnachten keine Risikobegnung mehr. Meine Freunde auch nicht.[NEWLINE]Und ich fahre täglich mit der Bahn quer du
1352934211641077760 : Schon intressant, wie viele Menschen aufgrund eines triftigen Anlasses an einem Samstag Mittag mit der S-Bahn unterwegs sind. #CoronaVirusDE #lockdown2021 #Berlin
1361944955103698944 : Crazy, in der Berliner U-Bahn ist die #Covid Pandemie schon vorbei. Unglaublich wie voll die ist und wie viele Menschen wieder zur Arbeit pendeln... #berlin
1363851241873428480 : Es ist Montag, 15 Uhr, die Berliner U-Bahn ist VOLL. So voll wie vor Corona. Man kann keinen Abstand halten.[NEWLINE][NEWLINE]Wie kann das „harter“ Lockdown sein? Wie s
1372453498470723584 : Covid19 Varianten verbreiten sich. Die S Bahn Berlin hilft fleissig mit und hat Fahrkarten kontrollen wieder aufgenommen, damits schneller geht #CoronaVirusDE
1374750711301484547 : Ich habe die Arbeit und die Kollegen während der Quarantäne vermisst.[NEWLINE]Andere Dinge wieder gar nicht...[NEWLINE]Penisnasen oder Maskenlosen in Bahn und Bahnhof
1375159963094552583 : Pandemie[NEWLINE]Ein Reisender kommt aus Spanien, sein Flug nach Berlin wurde Teil Storniert, dieser endet in Dortmund. Von da aus geht's in die U-Bahn, Fernzug mit Ums
1376566683855376388 : An April sollen in Berlin geflüchtete geimpft werden! #Impfdesaster #NieMehrCDU #Coronakrise [NEWLINE]Macht total Sinn! Die die das Geld verdienen und mit Bus und Bahn.
1377326647146393607 : Könnt ihr euch noch an Berlin zu Ostern erinnern? So mit tausenden Touristen? Die ihren Rucksack auf dem Bauch tragen, weil sie Angst haben, in der U-Bahn ausgeraubt zu
Enter your search query: Malaria covid
1344789923556175878 : i was in and out of over 100 different locations in the space of 3 months during covid lockdown buh God no let common malaria see us sef..
1345296698722148352 : The goal is to prevent the rise of CoVID-19 to reach levels of malaria and HIV related deaths. But no "you are giving SARS-CoV-2 too much
1345383461574610946 : Nick Searcy had the best line in a detective series that I think I have ever heard. He was describing a criminal they were pursuing. "He's
1345412318201401344 : "Tuberculosis still kills 1.5 million people a year, AIDS 700,000, malaria 400,000, and so on, and we barely bat an eye. When COVID-19 sto
1345452565908090883 : YT people thinking they're resistant to Covid like malaria didn't almost wipe them out.
1345463941154172928 : Me seeing people testing positive for COVID and I'm fretting cos the malaria no gree go☺
1345477945234747393 : If globalization will be criticized for the introduction of covid-19 pandemic to us, then what should take the blame of malaria disease th
1345509224701440004 : Why is COVID-19 an epidemic[NEWLINE]- cause one should not proscribe antibiotics against regular cold[NEWLINE]- cause you should not put p
1345813541530566659 : Has anybody noticed that Widal results comes out highly reactive in patients with Covid 19? So some people tell them it's Malaria/typhoid
1345820835764973572 : Malaria is more deadly than COVID-19[NEWLINE][NEWLINE]How many times did this ring in my head?[NEWLINE]#loveworld[NEWLINE]#PastorChris [NE
1345832852026494979 : A lot of people actually have covid but they will be disguising as malaria ☺
1345843494829645825 : Ebola & Malaria were all eradicated through science if vaccination can save us from current affliction we're suffering, this deadly harass
1345861815868268547 : People have been asking why Covid does not seem to be such a big issue in many African countries. My hypothesis is that in many countries
1345872180106702848 : Nigerians just grew huge pair of balls and decided to be throwing parties like covid 19 went on recess. Those of you on the TL claiming to
1346089360521887744 : Today's mini-shift[NEWLINE]Severe COVID req ICU[NEWLINE]Severe COVID req ICU[NEWLINE]Fracture mid-shaft femur[NEWLINE]Non traumatic SAH se
1346198670757150720 : The world should not receive COVID-19 vaccines from China. It's like asking the mosquito to treat malaria. ☺
1346199914179657731 : People that died from COVID died from other conditions and were told to label them COVID deaths. COVID is survivable under normal circumst
134621892989888647 : Please the malaria that you think you have may not be malaria ooo! Go and get a Covid test. [NEWLINE][NEWLINE]Thanks.
1346219318418268171 : The vast majority of people in Lagos will either get the new strain of Covid or the revamped strange Malaria
1346224805717159941 : Hospitals are filled to the brim and you don't believe Covid is real?[NEWLINE]So everyone somehow has Malaria + cough?[NEWLINE]Isokay.
Enter your search query: mal
```

- The information need „show me tweets of people who talk about the side effects of malaria and COVID vaccines“ is a phrase query
- We need **3 or 4 words** to describe the information need → e.g. **malaria covid vaccine effect**
- However our query engine is currently build to allow phrase queries only for 2 words, therefore the query „**malaria covid**“ is the best we can do with this implementation.
- But it is easy to extend the query engine to allow phrase queries with more than 2 words, by consecutively intersecting the postings lists of the words.
 - E.g. $r_1 = \text{query}(\text{term}_1, \text{term}_2)$
 - Intersect r_1 with $\text{query}(\text{term}_3)$ etc.